

New FindersKeepers v2 Frontend Design (UI/UX Overhaul)

Design Goals & Visual Style

The redesigned **FindersKeepers v2** frontend will emphasize real-time feedback, clarity, and modern aesthetics. It will incorporate a sleek **purple and teal** color scheme with a dark theme, creating a futuristic feel that matches the “AI GOD MODE” concept ¹. The current interface is already professional (built with Material-UI and a proper navigation structure ²), but the new design will modernize this with **Tailwind CSS** and **ShadCN UI** components for a fully custom look. All animations will be **tasteful yet flashy** – using **Framer Motion** for smooth transitions and subtle glowing highlights in purple/teal to draw attention without overwhelming the user. The goal is to maintain a polished, **production-ready UI** while improving functional feedback and aligning with the project’s cutting-edge nature.

Key design objectives include:

- **Real-Time Insight:** Expose the inner workings of the AI MCP server (session management, tool use, memory logging) in real-time, with clear indicators for successes/failures. The MCP server provides persistent session memory and intercepts all conversation and tool events ³, so the UI should surface this by showing live events and status changes.
- **Modern Tech Stack:** Use **React + Vite** (existing setup) with **Tailwind CSS** for styling and **ShadCN/ Radix UI** for accessible components (buttons, menus, cards, etc.). **Recharts** will render charts for usage stats, and **Framer Motion** will handle micro-interactions (e.g. slide-in panels, fade-in updates).
- **Purple & Teal Theme:** Define custom Tailwind theme colors (e.g. a rich purple for primary elements, and a bright teal for accents). These colors will be used for highlights, interactive states, and possibly a subtle gradient background. The overall look should be sleek and “techy” – for example, teal highlight glows on active elements or purple neon outlines on focus states.
- **Responsive & Intuitive:** Layouts will be responsive, adapting from widescreen dashboards to smaller windows. Use concise icons and headings so users can **scan status at a glance**. Short paragraphs of text and tooltips will explain any complex info (keeping text blocks small for readability).
- **Performance Toggle:** Keep resource usage in mind – heavy real-time features (like continuously updating logs or animations) will be **disabled by default** and only enabled when the user opts in. This ensures the UI remains light unless the user explicitly needs detailed live monitoring ⁴. We will utilize context and settings to control subscriptions to live data, and leverage efficient rendering (e.g. virtualization for long log lists).

Overall Layout & Navigation

The application will be organized into clear sections with a consistent navigation. A sidebar or top navigation bar will list the main pages of the app (with iconography for quick recognition). Based on the existing routes ⁵, the primary sections are: **Dashboard, Sessions, Tools Log, Monitoring, Settings** (plus

existing pages for Chat, Documents, Knowledge Graph, etc., which will be visually aligned with the new theme). Navigation will use Tailwind + ShadCN components (for example, a collapsible sidebar with Radix UI's navigation menu or a set of ShadCN `Button` links). Each page's title will be clearly shown at the top, with breadcrumb or section info if needed.

- **Sidebar Navigation:** A vertical sidebar (with a dark purple background) can house the section links. Each link will have a corresponding icon (e.g. a home/dashboard icon for Dashboard, a list icon for Sessions, a terminal icon for Logs, a heartbeat monitor icon for Monitoring, and a gear icon for Settings). The active page's link will be highlighted in teal. This persistent sidebar makes it easy to jump between sections that monitor different aspects of the system.
- **Header Bar:** A top header will display context info and quick controls. On the right side, it might show the current **Active Agent** (with an icon or avatar and name) and a status dot indicating if it's currently online/engaged. For example, if the Claude Desktop session is active, it could show "Active Agent: Claude (Desktop)" with a small purple glow, and if a VSCode-based agent session is active instead, the name would update accordingly. This gives a constant cue of **which agent is currently in focus or generating output**. The header also can include a "Live Monitor" toggle button (e.g. a stylish toggle or switch that enables the real-time capture console) and a shortcut to Settings.
- **Content Area:** The rest of the screen (to the right of the sidebar, below the header) is the page content. This area will be scrollable and padded. We will use Tailwind utility classes for grids and flex layouts to arrange content in each page as described below.

All pages will share the same theme and styles so the experience is cohesive. The navigation and header are fixed, and all icons, buttons, and text use the purple/teal accented theme. We will preserve the existing dark/light mode toggle functionality ⁶ ⁷ by implementing Tailwind's dark mode classes (or CSS variables) – allowing users to switch if needed (though the primary design is optimized for dark mode with neon accents).

Homepage Dashboard

The **Dashboard** is the high-level overview and entry point, giving the user a summary of system status and recent activity. It will be redesigned to be visually engaging and informative, combining text, icons, and charts in a clean layout. (Previously, a complex dashboard caused issues and was simplified to just an overview ⁴. Our design will reintroduce advanced features carefully, ensuring they load only when needed.)

Layout & Components on Dashboard:

- **Welcome/Project Header:** At the top, a heading like "FindersKeepers Overview" or a welcome message, possibly alongside the current date/time or user name. This can be a simple `<Typography>` heading styled in teal. Below it, a subheading or tagline can remind that this is "AI GOD MODE – Persistent AI Memory System" to set context.
- **Quick Stats Cards:** A grid of **statistic cards** summarizing key metrics:
 - **Total Sessions:** number of AI sessions recorded (with an icon, e.g. a SessionIcon, and number) ⁸.
 - **Active Sessions:** count of sessions currently active (with a green play icon) ⁹.
 - **Avg Session Duration:** average duration of sessions (with a clock icon, e.g. showing minutes) ¹⁰.
 - **Error Rate:** percentage of sessions that ended in error (with a warning icon) ¹¹.

These four metrics are already calculated in the system ¹² ¹³ ; we will present them as cards with large numbers and labels. Each card will use a Tailwind `card` style (or a ShadCN Card component) with a dark background and a colored left border or icon circle (purple, teal, or semantic colors for success/warning). This gives a quick system health snapshot (e.g., if error rate is >0, it will be highlighted in orange/red). - **Active Agent & Session State:** A panel showing which agent is currently active and the state of the “AI GOD MODE” session. For example: - *Active Agent:* “Claude Desktop” or “VSCode” or others, with an avatar or icon. This updates in real-time if a new session starts with a different agent. If multiple agents can be active, this panel could list each active agent with its status. - *Session Status:* If a session is ongoing, show “Session ACTIVE” in green; if the AI God Mode session was ended and summarized, show “Session ENDED – Summary available” with a link or button to view the summary. The UI knows when a session ends because `end_session()` in MCP produces a summary ¹⁴ , so we can check if the current session has a summary stored. A **Resume** button might also appear if the last session is ended (to allow quick resumption, since `resume_session()` loads previous context ¹⁴). This panel essentially reflects the session lifecycle (Active/Ended) in real-time. - **Recent Activity Feed (Mini):** A small recent log panel showing the last ~5 events in the system. This can be a condensed version of the live monitor – for example, the most recent tool invocation or message. Each line might say, e.g., “ *Search Tool* invoked by Claude – **Success**” or “⚠ *DB Insert* failed for conversation log”. This gives a quick glimpse of what’s happening. Each entry will have a timestamp and a colored icon (green check for success, red alert for failures, etc.). We will pull from the same event stream as the live monitor, but just display a few recent ones on the dashboard. (If real-time is turned off, this feed could instead show the last events fetched periodically, or remain static.) - **System Health Summary:** A section summarizing FastAPI and database statuses: - This might be a set of **service status badges** for the core components: FastAPI backend, PostgreSQL, Qdrant (Vector DB), Neo4j (Graph DB), Redis, Ollama (LLM). Each will show an icon (e.g. database icon, graph icon) and a small label like “FastAPI: Up” or “Neo4j: Down”. We can color these badges (green for up, red for down, yellow for degraded) using data from the `SystemHealth` API ¹⁵ . For example, if `systemHealth.services.fastapi.status = "up"` , show a green “FastAPI ✓” badge; if one is down, show it in red with an error icon. This mirrors the information in the backend’s health check, letting the user see at a glance if any service is offline. - Additionally, if desired, we could show the **last sync time** for conversation logs or summaries (e.g., “Last log saved at 01:23:45” if the system reports that). However, since reliable sync is mostly indicated by no errors, we might instead rely on the absence of alerts or an “All systems nominal” message. - **Usage Charts (Toggleable):** For a more dynamic overview, the dashboard can include one or two small charts (using Recharts) that update periodically or on demand: - For example, a line chart of “Messages processed per minute (last 10 minutes)” or “Vector searches vs Graph searches over time”. Another could be a pie chart or bar chart of “Tool usage distribution in the current session” (how many file reads, writes, searches, etc. have been done). Since the MCP tracks each action with a type ¹⁶ , we can aggregate that data for a session or a day. - By default, to save resources, we might show these charts with static data (e.g., summary of today’s usage). The Settings can allow “real-time charts” which, if enabled, subscribe to live data or refresh frequently. This addresses the roadmap goal of advanced real-time dashboard charts ¹⁷ while keeping performance optional. - **Animated Feedback:** Subtle animations will make the dashboard feel alive. For instance, if a new session starts, the “Active Sessions” card number will increment with a brief highlight or count-up effect. The service status badges might pulse or glow if a status changes (e.g., FastAPI goes down could flash the badge red briefly). Chart elements can animate when loaded (using Recharts’ built-in animations or Framer Motion’s `motion.div` around them for fade-in). These polished touches ensure the dashboard isn’t static, and changes in system state attract the user’s eye immediately.

Overall, the dashboard serves as a **mission control**: at a single glance, the user can see active agent/session state, whether all services are OK, and what the AI system is doing recently. It's visually engaging with the purple/teal theme – for example, cards might have a faint purple gradient background and teal numeric text, and icons can use the theme colors – yet it remains functional and not overly cluttered (each section is in its own panel or card, separated by whitespace).

Session Viewer (Active & Historic Sessions)

The **Sessions** page will allow the user to review all AI sessions (both currently active and past sessions) in detail. It builds on the existing Agent Sessions interface, which already displays session lists and stats ¹² ¹⁸. We will refine this with the new design and ensure it reflects real-time session management and persistent memory usage.

Layout & Features:

- **Session Stats Header:** At the top of the page, we'll show overall session statistics in small cards or badges (similar to the dashboard but focused on sessions):
- *Total Sessions, Active Now, Completed Today, Average Duration, Error Rate*, etc. (These were computed in the current code and displayed as four cards ¹² ¹³; we can include “completed today” as well since it's tracked.) Each stat will be a Tailwind-styled card with an icon and number, like on the dashboard, giving context for the list below.
- **Filters & Search:** Just above the session list, include controls to filter sessions:
 - A search box to filter by session ID or project name (as implemented) ¹⁹.
 - A dropdown (select) to filter by status (All, Active, Completed, Error, Terminated) ²⁰.
 - A dropdown to filter by agent type (All, Claude, GPT, Local, etc.) ²¹.

These will be re-created with ShadCN form components (e.g. `<Select>` with a list of options) styled in the new theme. Filtering happens instantly on the client side as currently done. The filters help the user quickly find a specific session or narrow the view.

- **Session List/Table:** The sessions will be displayed in a table or list view that is scrollable within the page:

- Columns: **Session ID**, **Agent** (type or name), **Status**, **Started (timestamp)**, **Duration**, **# Actions**, **Operations** (actions like view or terminate). This matches the existing columns ²². We will format these with Tailwind: e.g., a flex table or CSS grid.
- Each row will have visual cues:
- The **Session ID** can be monospaced and truncated (with a copy-to-clipboard icon on hover for convenience).
- **Agent** can be shown as a colored badge or chip (Claude, GPT, etc.) – e.g., a small pill with the agent name.
- **Status** will be a colored badge: Active (green), Completed (blue), Error (red), Terminated (gray). We'll use the status to choose the color (similar to how `getStatusColor` was used in MUI code) ²³.
- **Start time** will be formatted in a human-readable way (e.g. `toLocaleString` as current code does) ²⁴.
- **Duration** will display as “X min Y sec” or “–” if ongoing (the code's `formatDuration` does that) ²⁴.
- **# Actions** (total actions in that session) shows how many tool invocations or notable actions occurred. This gives a sense of how active the session was.
- **Operations (Actions column):** We'll include icon buttons for “View Details” (an eye icon) and if the session is active, an “End Session” (stop icon) ²⁵ ²⁶. We'll style these with ShadCN's `IconButton` or just Tailwind classes (small circular buttons with hover effects). The “End” button will be red and possibly require confirmation (as implemented).
- Each row could have a subtle animation on update – for example, if a session changes status from active to completed in real-time, the row could flash or an icon could appear indicating it ended. The system subscribes to session updates via WebSocket ²⁷, so if a session's status or action count changes, we'll update the row live (e.g., incrementing the action count, changing the status badge to “Completed” when it ends).
- Active sessions might be

highlighted at the top of the list or visually marked (perhaps a small green dot next to the ID or a subtle glow) to distinguish them. - **Session Detail View:** When the user clicks “View Details” on a session, a detailed view will be shown. This can be implemented as: - A **modal dialog** (as currently done with MUI Dialog ²⁸) that appears over the Sessions page, or - A dedicated route/section on the page that expands (e.g., an accordion or side panel). A modal is straightforward: we’ll use a ShadCN Dialog component for consistency. - In the **Session Details**, we show comprehensive information: - Basic info at top: Agent type, Session ID, Status, Start and End time, Duration. These can be displayed in a two-column grid of labels and values for clarity ²⁹ ³⁰. - If the session has ended and produced a **Summary**, we will show the summary text here. For example, “**Session Summary:** (some paragraphs of AI-generated insights)” stored from the end_session ¹⁴. This might be within a collapsible panel if it’s long, or simply a scrollable text area. The summary gives an overview of what happened, which is a key part of persistent memory across sessions. - A list of **Actions taken during the session:** This corresponds to the AgentAction list. We will list each action in chronological order (the backend provides them via `getSessionActions`). For each action, show: - An icon representing the action type (e.g. file icon for `file_read`/`file_write`, terminal for `command_execute`, cloud/API icon for `api_call`, search icon for `search`). The `action_type` field covers these categories ³¹. - A short description: we can use the action type name and maybe a key detail. For example, if `action_type` is `search`, the details might contain the query – we could display “Search: <query>” if available from `action.details`. If `action_type` is `file_read`, show which file was accessed (from `files_affected`). The data structure has a `details` field for extra info ³². Even if we just display the type and timestamp initially (as currently: “search – 2025-08-14 01:40:23”), we can enhance this. - A timestamp of when it occurred. - Success or failure indicator: since each AgentAction has a `success: boolean` and possibly an `error_message` ³³, we can show a ✓ or ⚠ icon colored appropriately. For example, a file write might fail, in which case we show a red error icon with a tooltip of the error message. These actions can be listed in a scrollable list inside the modal ³⁴. We’ll use a styled list (or even a table) with small rows for each action. Real-time note: if the session is active and ongoing while the user has the detail view open, new actions should live-update into this list (via the WebSocket subscription). We will append new actions to the list with a small highlight (maybe a brief teal background flash to draw attention). - If **no actions** recorded, we display an informative note (e.g., “No actions were recorded for this session” similar to the current alert) ³⁵. - Possibly include a button to **Resume** this session (if it’s ended) from the detail view, which would trigger the MCP resume on the backend and open the Chat interface with that context. This depends on whether resuming via UI is in scope; if not, we simply show info. - The modal will have a close button. It will be styled in line with the theme (dark background, maybe a purple border at top to emphasize it, etc.). By using a modal, the user can quickly inspect a session and close it to return to the list.

- **Live Updates & Toggles:** The session viewer will by default update in real-time for key changes (new session started, status changes) since these are relatively infrequent and important. However, if needed, we can respect a global “real-time updates” setting: for example, if the user disables real-time updates globally, the session list would not auto-update; instead the user would rely on manual refresh. A manual **Refresh** button is present (as in current UI) ³⁶ to fetch the latest sessions on demand. This gives control in case of performance issues or if the user wants to freeze the view for analysis.
- **Visual Enhancements:** We will use the purple/teal palette to make this page visually clear:
 - Status badges colored by status (as mentioned).
 - The filter section could have teal outlined inputs and a teal search icon for the search bar.
 - The table rows on hover might highlight with a subtle purple background.
 - Active session rows might have a small glowing border on the left (teal glow) to signal they’re active.

- **Animations:** When filters are applied, we can animate the list filtering (e.g., fade out and in). When a new session appears (user starts a session elsewhere), we could slide down a new row at the top with a highlight.

By providing both an overview of all sessions and drill-down capability, the Session Viewer makes the **persistent memory tracking** tangible: you can see every session's lifespan and what happened in it. This reflects the system's ability to maintain context across sessions and log everything ³. The user can easily identify if a session failed or was terminated (and see the error), reinforcing trust that no conversation or tool use goes unaccounted.

Tool Invocation Log Viewer

The **Tool Invocation Log** page (or "Actions Log") will present a real-time stream and history of all tool usage and other actions the agents have performed via the MCP. This is essentially a consolidated audit trail of AgentActions across sessions. It addresses the need to see *"which tools are being invoked through fk2-mcp"* and *"what messages are flowing through the MCP"* in one place.

Layout & Features:

- **Filter Controls:** Similar to the session list, we can allow filtering of the log:
- By **Agent/Session** – e.g., a dropdown to show actions from "All sessions" or narrow to a particular session or agent. We might list active session IDs or agent types.
- By **Action Type** – a multi-select or checklist for the types (`file_read`, `file_write`, `command_execute`, `api_call`, `search`). The user could, for example, filter to only see "search" actions to monitor knowledge queries.
- By **Status** – success or failure, if the user wants to see errors only.
- A text search could allow keyword filtering (e.g., file name or command snippet). These filters help manage the potentially large volume of log entries.
- **Log Entries List:** The main content is a chronological list of tool invocation events. This can be styled as an interactive list or table with one action per line (newest at top by default).
- Each entry will include key details:
 - **Timestamp** – when the action occurred (formatted in HH:MM:SS or full date if older). We might show relative time (e.g., "5 minutes ago" for recent entries, switching to date for older ones).
 - **Session/Agent** – which session or agent triggered it (could display the first 6-8 chars of session ID or the agent name). Possibly link this to the Session Details for context.
 - **Action Description** – a concise description of the action: *We derive this from* `action_type` *and* `details`. For example: "**Search** – query= `How to deploy?`" or "**File Write** – `/project/README.md`" or "**API Call** – OpenAI API". If details are complex, we keep it short, but provide a way to see more (like a hover tooltip or expandable row).
 - **Outcome** – success or failure. A green check icon for success, or a red error icon if `success=false` ³³. If failure, we can allow the row to expand or tooltip to show the `error_message` for debugging.
 - Possibly **Duration** – how long the action took (since `duration_ms` is logged) ³⁷. We can show this as "(123 ms)" or similar, perhaps in a subdued text after the description. This helps identify slow operations.

- Example entry: *"01:45:12 — Claude (Session 123abc...) — 🔍 Vector Search 'GPU acceleration' – Success (250ms)"*.
- The list will use a monospaced font for technical details to evoke a console feel, and use color accents for different parts (timestamp in neutral grey, agent in teal, action type bolded in purple, success in green or failure in red).
- **Icons:** Each action type can be preceded by a specific icon for quick visual parsing: e.g. a magnifying glass for search, a file icon for file actions, a terminal/command icon for executes, a cloud or link icon for API calls, etc. This makes the log easier to skim.
- **Live Streaming:** The log viewer will have the capability to display events in real-time. When enabled (via a toggle or if the global real-time setting is on), new tool invocations will appear in the list immediately as they happen. We will use the MCP's event stream (e.g., `action_completed` events in WebSocket) to populate this. When a new entry comes in:
 - It will be inserted at the top of the list (assuming descending chronological order).
 - We can animate its appearance (slide down or highlight) to draw attention.
 - If the user has scrolled up to view older events, we might not auto-scroll to bottom (to avoid disrupting them), but we can provide a "New events" notification or a button to scroll to latest.
- **Pagination/Length:** Over time, this log can grow large. We will implement either virtual scrolling (so rendering remains performant) or simple pagination (e.g., show last 50 by default, with a "Load more" button to fetch older entries). This prevents the UI from slowing down due to an extremely long list.
- **Manual Refresh & Pause:** If real-time is off or paused, the user can click a **Refresh** button to load recent log entries from the server (via an API call). Also, a **Pause** button can be provided to freeze the live updates (helpful if the user wants to inspect something without new lines jumping in).
- **Contextual Actions:** We might include small actions per log entry, such as:
 - Copy details to clipboard (for error messages or commands).
 - Jump to related session – clicking on the session ID or agent name could navigate to that session's detail view.
 - Expand for more info – if we have more data in `details` (like a full command or API payload), an expandable panel could show it indented beneath the entry.
- **Vector/Graph Search Feedback:** Since a key interest is tracking **vector search usage and fallback** and **graph search usage**, those will appear in this log:
 - We'll label vector searches distinctly (the action might come through as an `api_call` or `search` with details indicating Qdrant usage). We can parse details to see if it was a Qdrant query and label it **Vector Search**. If a vector search had no results and a fallback was triggered (e.g., perhaps calling a different endpoint), we could log an additional entry like "Vector search returned no results – falling back to keyword search" (possibly as a warning entry). These events could come from `system_alert` events or be inferred by the absence of results.
 - Similarly, graph queries to Neo4j can be labeled **Graph Query** in the log. If the graph endpoint is unreachable, we'd likely get an error action logged (which we'd show in red with the error message).
 - This way, the log viewer will explicitly show each attempt to use the knowledge bases (vector or graph), and whether it succeeded. For example: *"🔍 Vector Search – query 'X' – No results, fallback triggered"* could be one entry, followed by another entry for the fallback action (like a database search).
- **Visual Design:** The log page will have a somewhat "terminal" or developer console vibe, but dressed up with the theme colors and clean layout:

- The background of the log list could be a very dark color to contrast colored text (like a shade of near-black), possibly with slight transparency if over a gradient background, to give a heads-up display feel.
- Each log line might be separated by a subtle divider or alternating background stripe for readability.
- Hovering on a log entry could highlight it (e.g., light teal background) and show any action buttons (copy, expand).
- We'll use Tailwind to manage text sizes and colors (e.g., `text-sm text-gray-300` for timestamps, `text-teal-300` for agent, etc.).
- Important to maintain a balance: it should not look as dense as raw console output; adding spacing, icons, and coloring helps scan it easily.

By implementing the Tool Invocation Log, the user can **observe the AI's tool usage in real-time**, fulfilling the transparency requirement. It aligns with the MCP's design of capturing all actions and outcomes ³. If a FastAPI call fails or succeeds, the log will reflect that (e.g., an API call action with a success flag). Each entry's `success` field from the backend ³¹ will directly inform the UI's success/failure indication, so the user immediately knows if a tool execution (like a code run or web search) worked or hit an error. This page essentially acts as a diagnostics console for AI agent behavior.

FastAPI & Database Status Panel

Ensuring the backend and databases are in sync is critical, so we will provide a **Status Panel** focusing on FastAPI and the databases (and other services) to answer *"whether database insertions are occurring properly"* and show backend health. This can be presented either within the Dashboard (as a summary) or in the **System Monitoring** page for more detail. We propose a dedicated panel on the **Monitoring** page, with a mini status overview on the Dashboard.

Status Panel Features:

- **Service Health List:** We will list each relevant service with its status:
- **FastAPI Server** – e.g., *Up* or *Down*, plus maybe response time. Since the app can call a `/health` endpoint, we use that data. If `systemHealth.status` is overall healthy or degraded, we'll show that too ³⁸ ³⁹.
- **PostgreSQL (DB)** – status and maybe uptime or last error.
- **Qdrant (Vector DB)** – status (so the user knows vector search connectivity).
- **Neo4j (Graph DB)** – status (for graph queries).
- **Redis, Ollama LLM**, etc., as present.

We'll display these as a two-column grid or list: service name and a status indicator. Using the data structure from the backend, each service has a status field ¹⁵. We can augment with icons: e.g., a green check icon and "Online" label for up, a yellow warning icon "Degraded" if applicable, red error icon "Down" for not responding. If there's extra info like `uptime_percentage` or `last_error` ⁴⁰, we can surface that via a tooltip or expandable detail.

Example: PostgreSQL: ● Up (5ms response) – 24 docs, vectors synced.

Neo4j: ● Up – 37 nodes loaded.

Qdrant: ● Up – vector search OK.

The text in parentheses could come from stats (like number of docs/vectors, or the response time). - **FastAPI Insert Status:** To specifically address whether conversation logs and summaries are syncing: - We can track the last few database insert events for conversation logs. Possibly, each conversation message saving might emit an event (or at least errors would emit a `system_alert`). We will display a **“Conversation Log Sync”** status: e.g., “Last message logged at 01:44:10 – OK” or if an error happened, “Error logging message at 01:44:05 – see logs”. Essentially, if no errors have been flagged in the log, we assume it’s fine, and we show a green status. If an error occurred, we show a red status with the time and maybe an error snippet. - Similarly for **Session Summary** inserts: “Last session summary saved at [time]” once a session ends. If the latest ended session summary saved without issues, green check. - These could be small line items under FastAPI or DB, or a separate “Database Ops” section that lists recent operations and their status. - **Vector & Graph Ops:** Since vector search and graph search are critical, the panel can also monitor their usage: - For **Vector Search:** Show if the vector database is reachable (from the service status above). Also, track if queries are succeeding. For example, we might display “Vector Queries: 10 today, 100% success” or “Vector Queries: last failed at 13:20 (fallback used)”. This would require counting successes/failures which might be more detail than needed on the status panel; often, if Qdrant is “Up” and no alerts, one assumes queries succeeded. We can integrate this with the log: e.g., if any vector search had no result (fallback), perhaps mark the vector status with a small info icon indicating partial issues. - For **Graph Search:** similarly, if Neo4j is up, we assume queries work. If an issue (like a query threw an error), the system likely logs an alert. We could show “Graph Search: OK” or if errors: “Graph Search: Error in last query – see logs”. - **Database Sync Indicators:** Another useful element is showing if the various storages (Postgres, Qdrant, Neo4j) are in sync. The architecture uses a multi-modal storage approach ⁴¹, so we can add a quick check summary: - e.g., “Documents: 24 in Postgres, 2 vectors in Qdrant, 37 nodes in Neo4j ⁴²” (the roadmap confirms those counts in a success state). - If numbers are mismatched or if some pipeline is behind (like unprocessed documents), we can highlight that. The system monitoring data includes processing queue stats ⁴³; for instance, if `unprocessed_embeddings > 0`, it means some documents haven’t been vectorized yet. We could surface: “Embedding backlog: 3 documents pending” to inform the user that not everything is indexed. - These details might be more appropriate on a detailed Admin view, but we can at least indicate if everything is up-to-date or not. - **Manual Sync Controls (if needed):** On the Monitoring page, we might also include some admin controls (as seen in the AdminControlCenter code) like “Run vector reindex” or “Flush queues”. However, since the question is about design and not explicitly asking for admin actions, we can omit those from this answer. We’ll focus on display, not control, unless it’s simple like a “Refresh status” button. - **Auto-Refresh & Real-time:** The status panel will update either on a schedule (every X seconds polling an API) or via a WebSocket event (`health_update`) as seen in the code ⁴⁴. If real-time updates are enabled, we’ll subscribe to health updates and update the UI instantly when something changes (e.g., a service goes down). If disabled, the user can manually hit a refresh. We will indicate the “Last checked time” (the code stores `lastUpdate` for this purpose ⁴⁵) so the user knows how recent the info is. - **Visual Design:** This panel will use a simple and clear design: - Possibly a table-like layout where the first column is the service name and the second column is the status. We can use a light border or separator between items. - Service names can be in a muted purple text, statuses using colored badges or icons as described. - We will ensure this is easily readable at a glance (using icons and color coding primarily, with text for details). - On the Dashboard, we might only show a condensed version (e.g. only highlight if something is wrong, or just a count of all services OK). On the dedicated Monitoring page, we show full details.

By having this status panel, the user gains confidence that “FastAPI + database sync” is functioning. If, for example, FastAPI is up but an insertion to Postgres failed due to some issue, it would reflect either as FastAPI degraded (if it reports an error) or as a system alert for a failed DB insert – which would then be

visible in the live monitor. The UI will not hide such issues; the status panel could even have an **Alerts** section listing any recent errors (like “Failed to insert conversation at [time]”). This ties in with the log, but surfaced in a summary form.

In short, the Status Panel is the health dashboard for the system’s moving parts – giving green lights or red flags for each component and ensuring the user that persistent memory (which relies on those DB insertions) is being maintained properly. All of this reflects the backend’s direct integration approach (no silent failures) that the architecture achieved ³ ⁴⁶ .

Live Capture Monitor (Real-Time Console)

The **Live Capture Monitor** is a toggleable console view that shows a live feed of all important events flowing through the MCP system. This addresses the need to see “*what messages are flowing through the MCP and being logged*”, along with tool invocations, session events, and even system alerts, in real time. It’s essentially a developer/debug mode view – extremely useful for monitoring the AI God Mode in action.

Key characteristics:

- **Toggle Activation:** The monitor will **not** be open by default (to save resources and avoid distracting new users). It can be enabled in two ways:
 - Via a **persistent toggle button** in the header (for example, a “live feed” icon that glows when on). Clicking it slides the console in/out.
 - Through the **Settings** page, a switch for “Enable Live Monitor” (which could auto-toggle the header button as well).

The monitor could appear as a **dock** at the bottom of the screen (like a developer console). For instance, clicking the button slides up a pane from the bottom taking maybe 25% of screen height (resizable by dragging perhaps). - **Content of Monitor:** This console will display a chronological log of events, similar to the Tool Log but more comprehensive. Types of events to include: - **Messages:** User and AI messages as they go through MCP. E.g., “ User: <user’s prompt>” and “ AI: <assistant’s reply>”. We can possibly truncate long messages or allow expanding them, to keep the console manageable. These let the user verify that conversation messages are being captured and relayed. - **Tool Invocations:** (These will duplicate entries from the Tools Log, but here they’re in context with everything else). E.g., “ Tool invoked: *search – query=...* – Success.” - **Session Events:** “⚡ Session START (ID ..., Agent=Claude Desktop)” when a session begins, “ Session END (ID ..., duration X)” when ended, and “↔ Session RESUME (ID ...)” when resuming. These are key milestones the MCP server emits ⁴⁷ . - **Backend Acknowledgments:** If the MCP sends something to FastAPI, we might log “➡ Sent to FastAPI: /api/mcp/action – Status 200 OK” or if failed “Status 500 Error”. Essentially, whether each message or action was successfully processed by the FastAPI endpoints (this directly addresses “*whether they were successfully sent to FastAPI*”). The system likely logs accomplishment/failure as part of tracking ³ ; we can tie into that. For example, if a JSON-RPC call returns an error, we show it here. - **Database Inserts:** We can output events like “ DB Insert: Conversation log saved (Session XYZ)” or “⚠ DB Insert Failed: Error ...”. These would likely come through as system alerts from the backend if any failure occurs. Otherwise, we might not log every single successful insert (too verbose), but perhaps a periodic confirmation or just rely on the absence of errors. We can at least show when a session summary is saved: “ Session summary for Session XYZ stored.” - **Vector/Graph Operations:** “🔍 Vector search query for ‘ABC’ (result: 2 matches)” or “ Graph query for entity ‘XYZ’ (result: success)”. If a fallback occurs, it might show “Vector search had no results, defaulting to keyword search” as an info line. - **System Alerts / Errors:**

Any system-wide alerts (e.g., “ Neo4j connection lost” or “ Low GPU memory” if such events exist) will be shown prominently, likely in red with a siren icon. This ensures the user notices critical issues in the live feed.

Essentially, the live monitor is the union of all relevant event streams (conversation, actions, health alerts), giving a timeline of everything the system is doing. - **Formatting & UI:** The console will use a mono-space or semi-mono font for the log to distinguish it from normal UI text. Each line will be time-stamped. We might use a color stripe or label for the category of event: - For instance, message events could have a purple bar on the left, tool events teal, system alerts red, etc., or simply an icon prefix as listed. - We will indent or slightly offset AI vs user messages for readability (like chat bubbles in a minimal form, or just label them). - Use of **color**: success events in green text, errors in red text, info in neutral. For example, “200 OK” responses in green, error codes in red. - The background of this panel likely black or very dark gray (like a terminal). We can overlay text with slight transparency to give a HUD feeling. - It will be scrollable (with newest events at bottom or top – likely bottom like a terminal, so it autoscrolls down). - **Auto-Scroll & Controls:** By default, when open, it will auto-scroll as new events come (like a live tail of logs). The user can scroll up to pause auto-scroll. We will include a “Scroll to bottom” button that appears when scrolled up and new events arrive (common pattern in chat windows). - **Performance Consideration:** We will instantiate the WebSocket/event subscription **only when the console is enabled**. This prevents unnecessary load. The code’s structure (subscribe to events) allows dynamic on/off. If the console is closed, we unsubscribe from those detailed events (some basic ones like session_update might remain for other pages, but heavy message streaming can be toggled). - **Filtering:** Within the console, perhaps allow the user to filter out certain event types if the flow is too noisy. For instance, checkboxes for “Show messages / tool usage / system events”. This way if someone only cares about tool invocations, they could hide the raw conversation lines. - **Use Case:** Imagine the user is running an AI session in Claude Desktop; with the live monitor open, they can **literally watch each step**: the user prompt appears, the AI’s response appears, the AI decides to use a tool (a tool invocation event line shows up), the tool’s result maybe triggers an insight (logged), etc. It’s like watching the brain and memory of the AI in action, which is the essence of AI God Mode’s transparency. This UI gives immediate feedback on whether everything is flowing correctly – for example, if a message didn’t log, it would be obvious by a missing entry or an error line. - **Example snippet in use:** (User enables Live Monitor) - 01:45:14 [Session] ✂ Session START (ID: 12345678, Agent: Claude Desktop) - 01:45:15 [Message] User: "Find the deployment steps for project X." - 01:45:15 [Action] Tool invoked: searchVector(query="deployment steps project X") - 01:45:15 [Action] Tool result: 3 relevant documents found (Vector Search) - 01:45:16 [Message] AI: "According to the documentation, the steps are..." - 01:45:17 [FastAPI] ➡ Sent conversation log to API - 200 OK ³ - (... and so on ...) - In this snippet, you see session start, the user query, the AI tool use (vector search success), the AI’s answer, and the confirmation that FastAPI logged the message successfully. All of that would be visible to the user instantly.

If any part failed, e.g. vector search returned nothing and fell back, we might see: - 01:45:15 [Action] ⚠ Tool result: no matches, falling back to keyword search -

```
01:45:15 [Action] Tool invoked: searchDatabase(query="deployment steps project X") (So the user sees the fallback logic happening.)
```

- **Closing the Monitor:** The user can hide it by clicking the toggle again. When hidden, the component will stop receiving updates (to free resources). We'll ensure the state (like filtered types or scroll position) can reset or persist as appropriate.

The Live Monitor truly brings all the **persistent memory tracking and real-time tool invocation** to the forefront of the UI. It's especially useful for power users or during development/debugging of the AI's behavior. By making it toggleable, we keep it out of the way for casual use but readily available for deep inspection. This fulfills the requirement of showing the internal MCP message flow and tool usage live, complementing the higher-level views (which might only show aggregated or status info). It leverages the MCP's comprehensive conversation capture ³ and ensures **nothing happening in the background is hidden from the user**.

Settings & Configuration Page

The **Settings** page will provide user controls to configure the frontend's behavior, particularly toggling real-time features and adjusting the UI to their preferences. It will also adopt the new design style and improve on the current minimal settings page.

Settings Page Layout:

- **Appearance Settings:** The theme toggle (light/dark mode) will be presented prominently. We'll have a section "Appearance" with:
 - A toggle or button group for Light vs Dark mode (similar to current, but using ShadCN Toggle or Radix Switch). It will reflect the current mode and allow change. This persists to localStorage as implemented ⁴⁸ ⁴⁹.
 - Possibly an option to choose accent color presets (if we allow customizing the purple/teal scheme), though not required. We might simply stick with the defined palette.
- **Real-Time Feature Toggles:** A section "Live Updates & Monitoring":
 - **Global Real-Time Updates:** A master switch to enable/disable live updates across the app. If turned off, the app will not auto-refresh data or subscribe to websockets; the user would manually refresh session lists, etc. If on, real-time subscriptions (sessions, health, etc.) are active. This was indicated in current settings UI as a static "Enabled" label ⁵⁰ – we will make it an actual toggle. When this is off, the Live Monitor and perhaps tool log auto-stream should be unavailable (or require manual refresh).
 - **Live Monitor Console:** A toggle to enable the live capture console. This might do the same as the header button – perhaps if you enable it here, it auto-opens or at least "arms" the console to turn on for new sessions. However, since we have a direct UI control for it, this setting could simply be descriptive or an alternate way. We could integrate it such that if this is off, even clicking the header button won't open the console (for users who never want it). Probably better is to just manage via the header button, and use the global real-time toggle to determine if that feature can function.
 - **Session Auto-Refresh:** (If we want granular control) a toggle to auto-update the session list in real-time. But this overlaps with global real-time. Maybe not needed separately.

- **Real-Time Charts:** A toggle to enable live updating charts on the Dashboard/Monitoring. If off, charts update only on page load or manual trigger. If on, they might update every few seconds or via streaming data if available.
- Essentially, this section ensures the user can dial down the live features if they are running on limited resources, etc.
- **Feature Toggles:** A section "Features" could list major features (vector search, knowledge graph integration, etc.) with their status:
- The app's config might allow disabling vector search or graph features ⁵¹. For instance, if `features.vectorSearch` is false, maybe hide the Vector Search page. In settings, show a toggle "Vector Search Enabled" ⁵², which currently is just labeled enabled in the UI. We can actually allow the user to turn it off if needed (though if it's tied to backend capabilities, toggling might be more of an admin thing).
- "Knowledge Graph Enabled" – similarly. This way if a user doesn't want those sections or to reduce system load, they could disable them.
- "Session Monitoring Enabled" – which might correspond to the session management features (though likely always on because that's core).
- These toggles could just reflect the config (if not meant to be user-changed, we might just display them as indicators). But as a UI/UX improvement, letting advanced users configure which modules are active could be valuable.
- **Notifications/Alerts:** (If applicable) If we have any user notifications (for example, a browser notification when a session ends or an alert if a service goes down), settings to control those can be included. Not explicitly asked, but worth considering for completeness.
- **Save & Feedback:** Most settings will apply instantly (especially toggles) since it's frontend behavior. But if some require a restart or heavy change, we'll note it. The page will be kept simple: a list of labeled toggles in categories, using clear language (non-technical where possible, with perhaps a short description under each toggle).
- For example: *"Real-Time Updates – When enabled, the app will live-update session lists, logs, and status panels without manual refresh."*
- Under that, the toggle component. If disabled, we might even auto-hide the live monitor and pause websockets.

Visually, the Settings page will use the Card or Panel style similar to current (boxes with a border and heading) ⁵³ ⁵⁴, but with Tailwind styling: - Each section (Appearance, Live Updates, Features, etc.) can be a `<div>` with a border or just separated by headings. - Use the theme colors for accents (e.g., the toggle switches in teal when on). - It should remain clean and not too flashy, since this is a static config area – but it will consistently use our font and color scheme.

By providing these controls, we empower users to tailor the UI's level of feedback. For instance, a user on a weaker machine might turn off live monitor and charts to conserve performance, whereas a developer would turn everything on. The **Settings page** thus complements the dynamic monitoring features with user governance.

(Note: The existing settings UI in the project already lists Real-time updates, Vector search, Knowledge graph as "Enabled" indicators ⁵⁵ ⁵⁶. We will convert those into actual interactive toggles or at least reflect true/false from a config. If some features shouldn't be user-toggleable, we can still display their status, but the question suggests a settings panel to enable/disable real-time views, so we'll definitely implement toggling for those.)

Component Structure & Integration

To fit naturally into the `fk2_frontend` codebase (Vite + React), we will organize the new UI into reusable components and follow the project's patterns:

- **Theming:** We'll integrate Tailwind CSS (with a custom config for our colors) and possibly use ShadCN's pre-built components by copying them into the project (ShadCN provides Tailwind-compatible React components). This replaces the Material-UI theme usage. We'll remove MUI providers and instead wrap the app in perhaps a context for theme (Tailwind's dark mode can be toggled by adding a `dark` class to `html`).
- **Layout Components:** Create a `Sidebar` component for nav, a `Header` component for the top bar (showing active agent, etc.), and a `MainLayout` that composes them. These can reside in `src/components/layout`. The `Sidebar` will map to the Routes (Dashboard, Sessions, etc.), using Link components. The `Header` will contain the agent status and live monitor toggle.
- **Dashboard Components:** E.g. `DashboardStats` (for the quick stat cards), `ServiceStatusPanel`, and chart components (could be encapsulated, like `UsageChart`). If charts are similar, we might have one generic chart component configured via props for different data.
- **Sessions Page Components:** We can break it into `SessionStatsCards`, `SessionFilters`, `SessionTable`, and `SessionDetailsModal`. This keeps the page file lean. The table rows and modal can be their own components for reuse or clarity.
- **Log Components:** Perhaps a generic `LogTable` or `LogList` component that both the Tool Log page and Live Monitor might reuse to format events. Or treat them separately: `ToolLogTable` (with filter UI and list) and `LiveConsole` (the sliding console UI). Given similar data, we might share a smaller component like `LogEntry` to render an event line given its data, so both the console and log viewer ensure consistency in how an action or message is displayed.
- **Status Components:** A `ServiceStatusList` component will iterate through services and render each with appropriate styling (using data from context or props). We might also have a `SyncStatus` component if we break out database sync info separately.
- **Settings Components:** Use simple form components (we can utilize Radix UI Switch or Checkbox for toggles). Possibly each toggle could be its own small component bound to some global config state.
- **State Management:** The project appears to rely on context/hooks for WebSocket events (`useWebSocket` hook) and direct API calls for data. We will continue this:
 - Use the `useWebSocket` hook to subscribe/unsubscribe to channels for sessions, health, actions, etc., within our components. For example, `LiveConsole` subscribes to multiple event types.
 - A central store (could be Context or even just a lightweight Zustand store) might hold things like active agent info, list of sessions, etc., so that the Header can display the active agent easily (perhaps derive it as the agent of the active session).
 - We will ensure to **unsubscribe** in `useEffect` cleanup for any live feeds when components unmount or when toggles off, to avoid memory leaks or unnecessary traffic.
- **Performance Tuning:** Where possible, use React lazy loading for heavy pages (maybe the Monitoring page with charts) so it doesn't load until needed. Also use React's memoization for lists (e.g. if we have a large log list, ensure each entry has a stable key and perhaps use `React.memo` on the entry component to prevent re-renders when not needed).

- **Icons & Assets:** We can utilize a modern icon set (Heroicons or Lucide icons, or continue using Material Icons via icon fonts). ShadCN typically uses Lucide icons by default. We will pick icons that match our theme (perhaps outlined neon-style icons).
- **Consistency:** All pages will share common UI elements (the same card styles, the same badge styles for statuses, etc.). We might create a small design system within the project: e.g., a `<Badge variant="success">` component for statuses (styled with Tailwind classes for green, etc.), so that Session status, Service status, etc., all use the same Badge component with just different variants. This makes it easier to maintain and ensures the purple/teal scheme is applied uniformly.

In terms of **integration into the file structure**, these components will be placed under `src/components` and pages under `src/pages`, just like the existing ones (Dashboard, AgentSessions, etc.). We will replace or heavily refactor the existing page files (which currently use MUI) with our new implementations using Tailwind. This can be done incrementally, but given the new design, it might be a full replacement of the JSX structure in many cases (while preserving the underlying logic and API calls).

Finally, we make sure to test the new UI thoroughly with actual data: opening a session and verifying the live logs show up, ending a session updates the UI state, etc. The aim is to achieve the *"comprehensive GUI interface with advanced features"* as outlined in the roadmap ⁵⁷ – including the advanced dashboard, complete session management integration, and real-time monitoring – all wrapped in a modern, shiny new interface.

Conclusion

This new frontend design will transform **FindersKeepers v2** into a visually striking and highly informative control center for the AI GOD MODE system. By leveraging a purple/teal modern theme, up-to-date React tooling (Tailwind, ShadCN UI, Recharts, Framer Motion), and carefully designed components, the UI will not only look polished but also give immediate, meaningful feedback about the system's operations.

Crucially, every aspect of the MCP server's capabilities – persistent session memory, tool usage, conversation capture, success/failure tracking – will be reflected in the interface in real time. The user will be able to monitor which agent is active, what tools the AI is invoking, and the health of each subsystem at a glance. And with the ability to toggle live views, the user remains in control of resource usage and can choose between a quiet overview or a detailed live debug view.

Overall, this design prioritizes **functionality and feedback** (making the system's invisible processes visible and understandable) while adding a layer of **modern design polish** to match the sophistication of the backend. It will make managing and interacting with the FindersKeepers v2 system a more intuitive, engaging, and empowering experience ³ ¹⁷.

¹ ³ ¹⁴ ⁴¹ ⁴⁶ ⁴⁷ finderskeepers-v2_architecture.md

https://github.com/cai76/finderskeepers-v2/blob/a74f7925451a859b62058f05c32a25b8015ecb0a/finderskeepers-v2_architecture.md

² ⁴ ⁵ ¹⁷ ⁴² ⁵⁷ ROADMAP.md

<https://github.com/cai76/finderskeepers-v2/blob/a74f7925451a859b62058f05c32a25b8015ecb0a/ROADMAP.md>

6 7 48 49 50 52 53 54 55 56 **App.tsx**

<https://github.com/cain76/finderskeepers-v2/blob/a74f7925451a859b62058f05c32a25b8015ecb0a/frontend/src/App.tsx>

8 9 10 11 12 13 18 19 20 21 22 23 24 25 26 27 28 29 30 34 35 36 **AgentSessions.tsx**

<https://github.com/cain76/finderskeepers-v2/blob/a74f7925451a859b62058f05c32a25b8015ecb0a/frontend/src/pages/AgentSessions.tsx>

15 16 31 32 33 37 40 51 **index.ts**

<https://github.com/cain76/finderskeepers-v2/blob/a74f7925451a859b62058f05c32a25b8015ecb0a/frontend/src/types/index.ts>

38 39 **AdminControlCenter_broken.tsx**

https://github.com/cain76/finderskeepers-v2/blob/a74f7925451a859b62058f05c32a25b8015ecb0a/frontend/src/pages/AdminControlCenter_broken.tsx

43 44 45 **SystemMonitoring.tsx.backup**

<https://github.com/cain76/finderskeepers-v2/blob/a74f7925451a859b62058f05c32a25b8015ecb0a/frontend/src/pages/SystemMonitoring.tsx.backup>